

Yining Hong

[✉ yhong3@andrew.cmu.edu](mailto:yhong3@andrew.cmu.edu) [🔗 hyn0027.github.io](https://github.com/hyn0027) [🔗 hyn0027](https://www.linkedin.com/in/hyn0027/) [🔗 hyn0027](https://orcid.org/0009-0002-6347-8222)

Education

Carnegie Mellon University <i>Ph.D. in Software Engineering</i>	<i>Aug. 2024 - May 2029 (Expected)</i>
Tsinghua University <i>B.Eng. in Computer Science and B.Ec. in Economics & Finance</i>	<i>Sept. 2020 - July 2024</i>

○ **CQPA:** 4.11/4.00

○ **Advisors:** Prof. Christian Kästner and Dr. Chris Timperley

○ **Research Interest:** Safety/security for software with ML/Agent components; AI for software safety/security

○ **CGPA:** 3.91/4.00 , **Rank:** 2/12

○ **Thesis (Honored):** A Reinforcement Learning Framework for Training and Testing in Asset Portfolio Optimization

○ **Advisors:** Prof. Maosong Sun and Prof. Junliang Xing

Internships

Software Engineer Intern <i>Advanced Micro Devices (AMD)</i>	<i>Beijing, China</i> <i>Jan. 2024 – June 2024</i>
○ Optimized FPGA-based deep learning compiler to improve runtime performance. ○ Applied operational research methods to optimize storage unit allocation.	
Software Engineer Intern <i>Kuangshi Technology (Megvii)</i>	<i>Beijing, China</i> <i>Sept. 2023 – Dec. 2023</i>
○ Integrated machine learning models into autonomous driving systems. ○ Designed and implemented guardrails for object trajectory prediction models.	
Student Research Intern <i>Carnegie Mellon University</i>	<i>Pittsburgh, PA</i> <i>June. 2023 – Sept. 2023</i>
○ Conducted research on semantic data slicing for capability testing in language models. ○ Co-advised by Prof. Christian Kästner and Prof. Sherry Tongshuang Wu. ○ Co-authored a paper accepted at ASE 2024.	

Publications

- [1] **From Hazard Identification to Controller Design: Proactive and LLM-Supported Safety Engineering for ML-Powered Systems**
Yining Hong, Christopher S. Timperley, Christian Kästner.
International Conference on AI Engineering (CAIN), 2025.
- [2] **What Is Wrong with My Model? Identifying Systematic Problems with Semantic Data Slicing**
Chenyang Yang, **Yining Hong**, Grace Lewis, Tongshuang Wu, Christian Kästner.
International Conference on Automated Software Engineering (ASE), 2024.
- [3] **Two Heads Are Better Than One: Exploiting Both Sequence and Graph Models in AMR-To-Text Generation**
Yining Hong, Fanchao Qi, Maosong Sun.

On-going Research

Guaranteeing safety/security for LLM-based Agents

- Agents connect with tools that could be malicious or vulnerable. We explore symbolic mechanisms, including API security practices, information flow analysis, and temporal logic, to provide safety/security guarantees.

Monitoring for risk management for software systems with ML components

- ML components in software introduce new risks, requiring new methods of monitoring. This interview study explores how practitioners monitor subtle risks and how interventions change their mental model.

Projects

Hazard Finder

github.com/hyn0027/Hazard-Finder ↗

- Developed an LLM-powered tool to assist safety engineers in performing hazard analysis (specifically STPA) for ML-enabled software and AI agentic systems, supporting systematic safety/security risk management.

OpenAI Chat Helper

github.com/hyn0027/OpenAIChatHelper ↗

- Developed and maintained a Python package for managing asynchronous OpenAI Chat API interactions with strong data typing, prompt templating, markdown formatting, and streamlined workflow support.

RL Portfolio Optimization (Bachelor's Thesis)

github.com/hyn0027/RL-Portfolio-Optimization ↗

- Developed a unified Python framework for implementing and testing (deep) reinforcement learning methods for financial asset portfolio optimization.

Relational Database Management System

github.com/hyn0027/DataBase-System ↗

- Developed a complete relational database management system in C++ from the ground up, reusing only ANTLR parser generator.

Skills

Programming Languages: C/C++, Python, JavaScript/TypeScript, Rust, Solidity

Safety & Security: Threat Modeling, STRIDE, STPA, FMEA, HAZOP, FTA

DevOps/MLOps: Docker, Git, MLflow, GitHub/GitLab CI/CD, Jenkins, Make/CMake

Monitoring: Grafana, Prometheus

Web Development: FastAPI, React, Next.js, Vue.js, HTML/CSS, Django

ML: MCP, PyTorch, NumPy, Pandas, Hugging Face

UI/Game Development: Qt, Cocos

Other: MySQL, L^AT_EX, Verilog/SystemVerilog, VHDL, Vivado, RISC-V Assembly

Awards and Honors

Outstanding Graduate , Computer Science Department, Tsinghua University

July 2024

Comprehensive Excellence Scholarship (top 10%) , Tsinghua University

Nov. 2023

Citadel Securities Scholarship , Citadel Securities

Dec. 2022

Comprehensive Excellence Scholarship (top 10%) , Tsinghua University

Nov. 2022

Academic Excellence Scholarship (top 30%) , Tsinghua University

Nov. 2021

Sports Excellence Scholarship , Tsinghua University

Nov. 2021

First Prize , National Olympiad in Informatics in Provinces

Nov. 2018

First Prize , National Olympiad in Informatics in Provinces

Nov. 2017

Teaching & Mentoring

Volunteer Teaching Assistant (Outstanding Level)

Tsinghua University

Beijing, China

Mar. 2022 - July 2024

- Offered drop-in tutoring for computer science and finance courses.

Aarya Doshi (mentoring)

Undergraduate student from Georgia Institute of Technology

May 2025- Aug. 2025

- Explored safety analysis and control for AI agents.

Abhishek Satpathy (mentoring)

Undergraduate student from University of Virginia

May 2025- Aug. 2025

- Explored threat modeling for ML-enabled systems.